

# XGBoost ML Algorithm

Gavan VanOver

Philip Waymeyer

Widnie Dorilas

Zheding Zhao



*dmlc*  
***XGBoost***

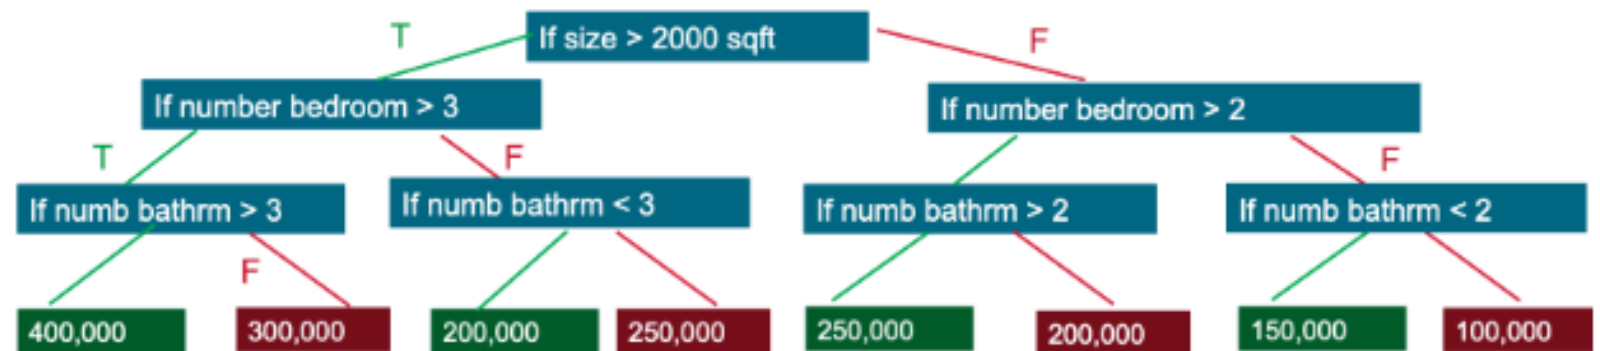
# Intro to XGBoost

- Essentially an optimized Gradient Boosting machine learning library
  - Gradient Boosting is already an advanced variation on Decision Trees
- Open-Source
- Compatible with C++, Java, Python, Julia, Perl, and Scala
- Garnered notoriety around 2014 for consistently outperforming other algorithms in competitive environments such as Kaggle competitions

# Intro to XGBoost

## Basic Algorithm

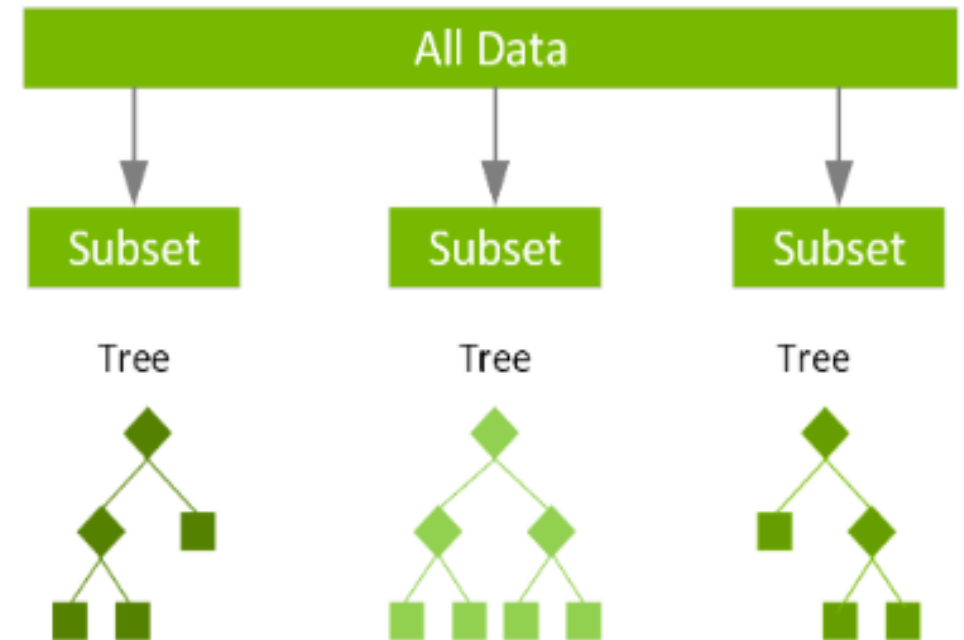
- Gradient Boost
  - Threshold Determination
  - Learning Rate
- Regularization
  - Lambda
  - Alpha
  - Gamma



# Intro to XGBoost


## Handle of Dataset & Beyond Statistics

- Approximate Greedy Algorithm
- Sparsity-Aware Split Finding
- Parallel Learning
- Weighted Quantile Sketch
- Cache-Aware Access
- Blocks for Out-of-Core Computation




# Data Processing for XGBoost

- First step is importing the data set you need
- Remove columns that will add noise
- XGBoost only manages numeric vectors
  - A way of changing categorical data to numeric data using *one-hot encoding*  
*AKA dummy contrast coding*
- *Null* values can be dealt with but don't have to be as XGBoost can internally handle missing values



# XGBoost Hyperparameters - General Parameters

- Most of the general parameters are managed by the algorithm, and they do not need to be manually set.
- A few that a user may want to change for one reason or another include:
- `booster[default=gbtree]`
  - Can be 'gbtree', 'gblinear', or 'dart'.
- `verbosity[default=1]`
  - Can be 0 (silent), 1 (warning), 2 (info), 3 (debug).
- `nthread[default to maximum number of threads available if not set]`




# XGBoost Hyperparameters - Booster Parameters

- XGBoost has two types of boosters: tree booster and linear booster.
- There are many booster parameters in XGBoost, so only a handful will be showcased.
- `eta[default=0.3]`
  - `Range[0,1]`
- `gamma[default=0], lambda[default=1], alpha[default=0]`
- `max_depth[default=6]`
- `scale_pos_weight[default=1]`
  - XGBoost recommends  $\text{sum}(\text{negative instances}) / \text{sum}(\text{positive instances})$ .
- `subsample[default=1]`



# XGBoost Hyperparameters - Learning Task Parameters

- These specify the learning task and the objective.
  - `objective[default=reg:squarederror]`
    - There are many values to use, such as `reg:logistic`, `binary:logistic`, `binary:hinge`, `count:poisson`, and `multi:softmax`.
  - `eval_metric[default according to objective]`
    - Defaults include `rmse` for regression, `error` for classification, and `mean average precision` for ranking. Multiple evaluation metrics can be added.
  - `seed[default=0]`
- 



# Advantages

- Works for both regression and classification
  - Rather than needing both linear and logistic regression
- Built-in ability to deal with missing values
- Nonparametric (no assumptions on data)
- Many hyperparameters to help the learning process
  - Regularization (prevent overfitting), learning rate, etc.
- Parallelization
  - Faster than other ensemble models (bagging or boosting)
  - Proficient memory usage

# Disadvantages

- If not handled properly the model is likely to overfit
- More difficult to interpret
  - Much more complex than linear or even logistic regression
- Extremely sensitive to outliers
- Doesn't perform well on unstructured data



Questions?

